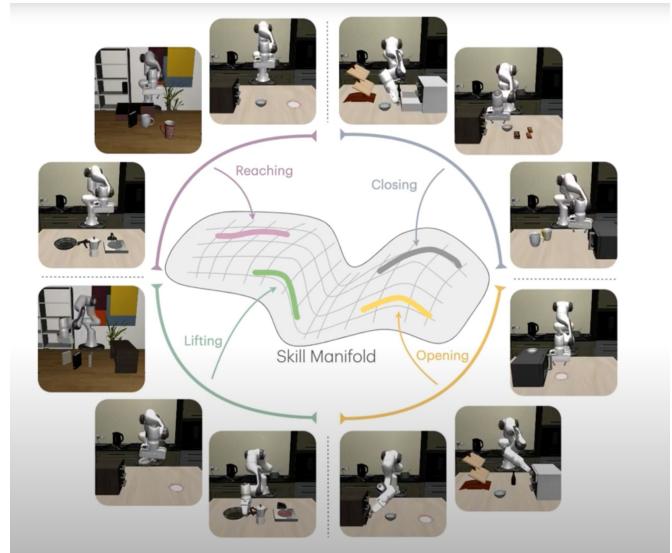
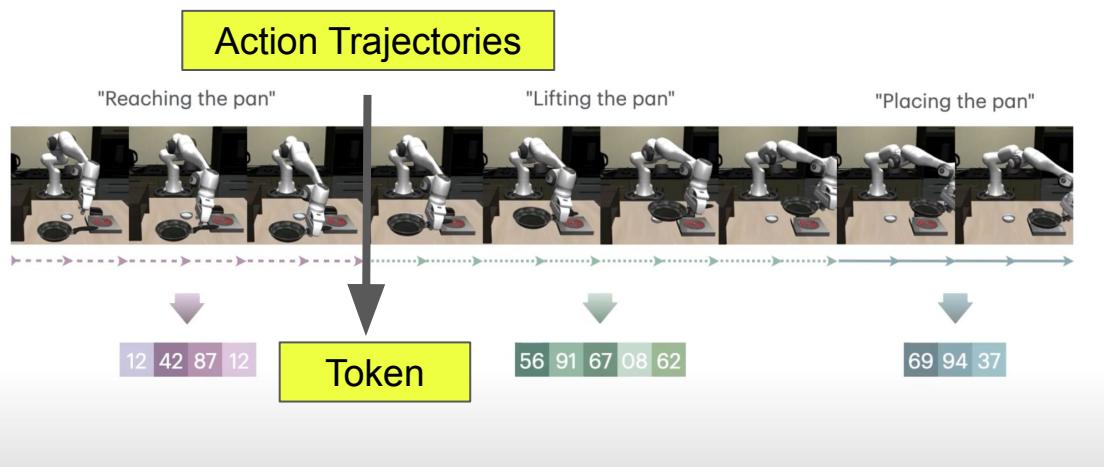


Skill-based Action Tokenizers

Pablo, Owen

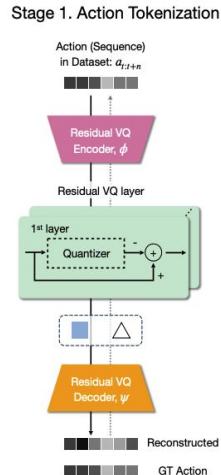
Research Objective

Can we compress a long demonstration into a short list of reusable skills?

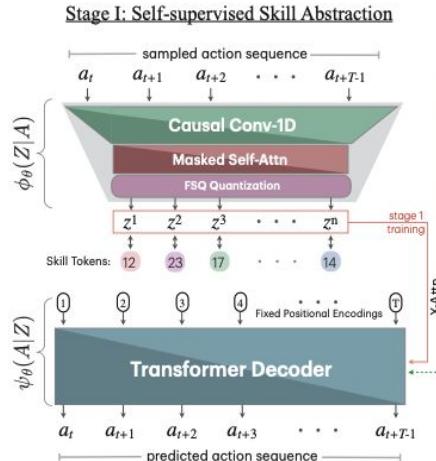


Related Works - Action Tokenizer

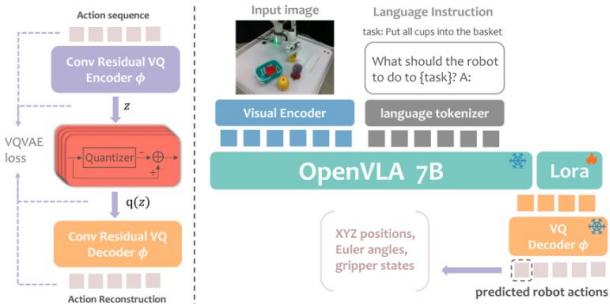
Vq-Bet



Quest



VQ-VLA



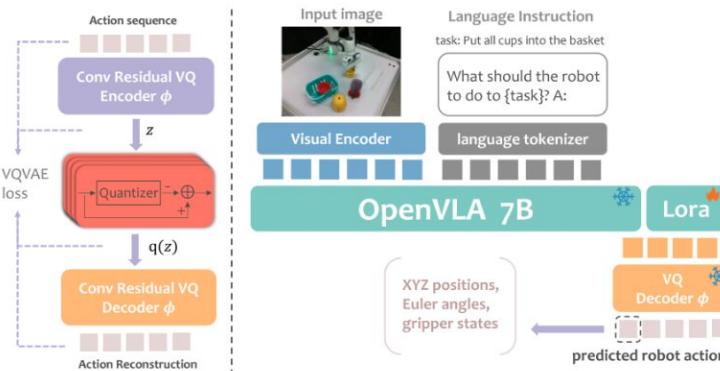
	VQ-BeT ('24)	QueST ('24)	VQ-VLA (ICCV, 25)
1. Compress : Raw snippet → token ID	<ul style="list-style-type: none"> Residual VQ-VAE: 2 codebooks jointly trained → up to 64–256 latent combos for each snippet. One token per chunk, but far richer code space. 	<ul style="list-style-type: none"> Conditional VAE with causal conv → produces a sequence of scalar-quantised codes (FSQ). (Each snippet can get 1–8 tokens, so variable-length skills.) 	VQ-VAE Decoder & VLA integration

Limitations of Related Approaches

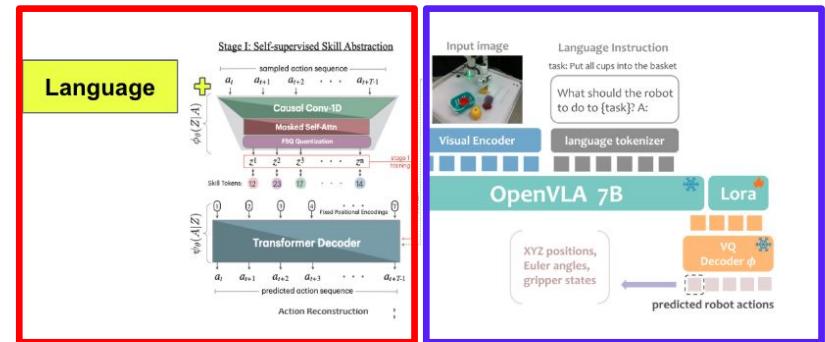
1. **Not conditioned:** Prior methods compress actions in isolation without considering task context or instructions
2. **Lack of proper metrics:** No quantitative way to evaluate if learned tokens are semantically meaningful, reusable

Main Idea

VQ-VLA

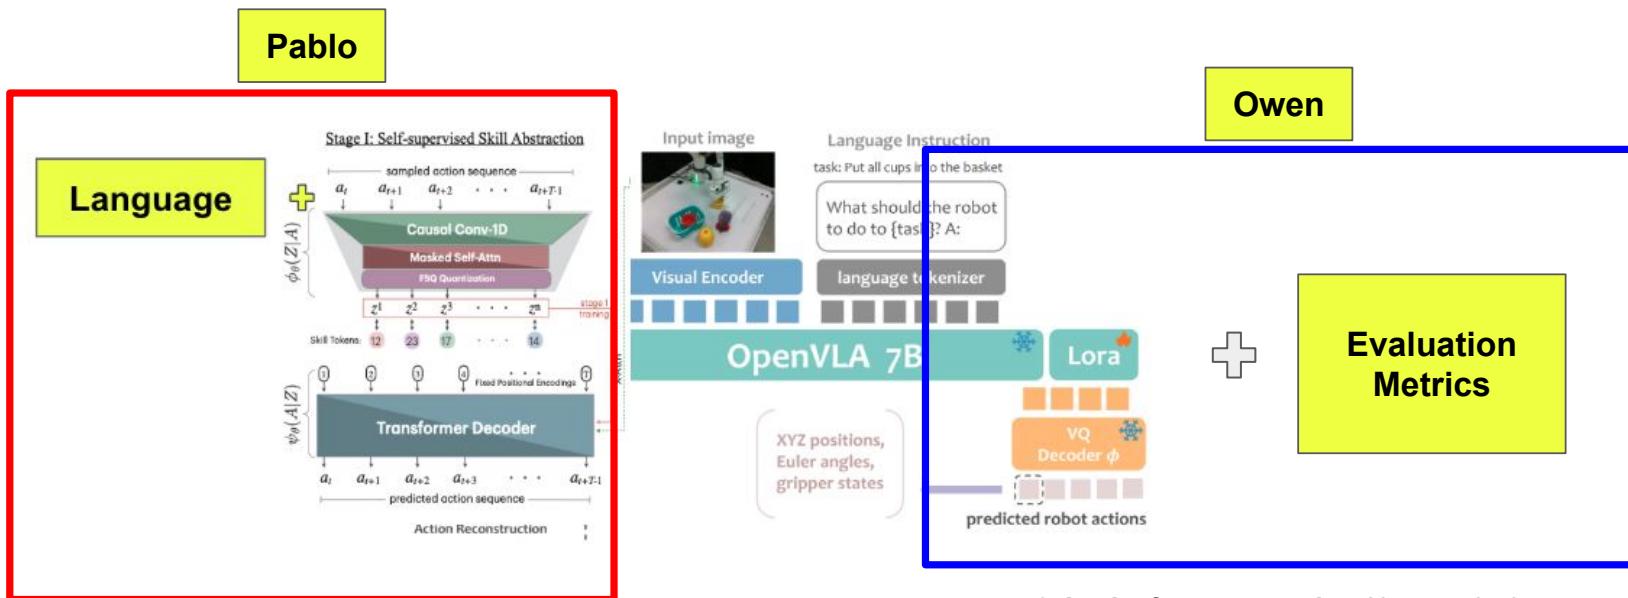


Ours



+ Quantitative Metric

How we are collaborating



1. Not conditioned: Prior methods compress actions in isolation without considering task context or instructions

2. Lack of proper metrics: No quantitative way to evaluate if learned tokens are semantically meaningful, reusable, or interpretable

+VLA Integration

Future Plans

~11

1. Research
2. Conditioned Architecture
3. Evaluation Metric
(BPA, InfoMEC, MDP-based)

11~

1. QueST-based Architecture
2. VLA Integration
3. Simulation (Libero)

~12

1. Simulation (Libero)
2. Real-robot

~1/26

1. Writing

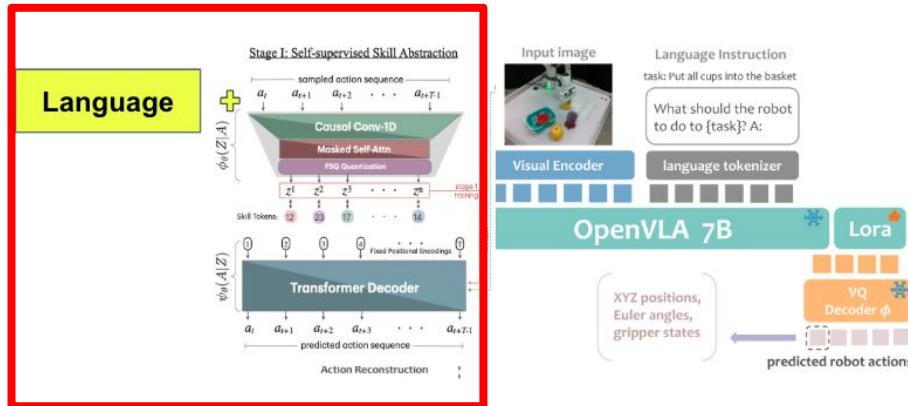
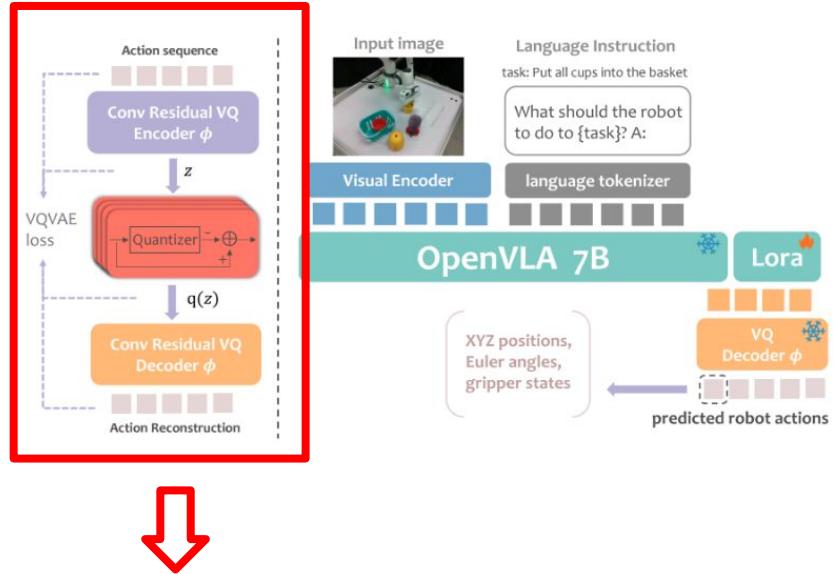
1. Conditioned Skill Action Tokenizer

Why this is needed?

Current VLA tokenizers mainly prioritize training efficiency but neglect long-horizon scalability and reusability.

Make action tokenizer to generate skill-based action not few action tokens.

> Encode entire motion primitives (e.g., “reach forward smoothly”) as single compressed tokens, enabling coherent multi-step execution.

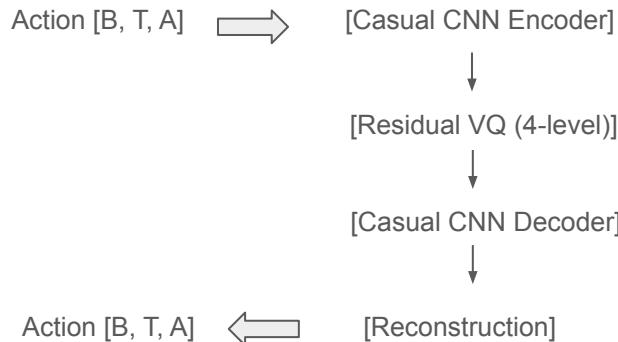


VQ-VLA: Improving Vision-Language-Action Models via Scaling Vector-Quantized Action Tokenizers

VQ-VLA

Step1

VQ-VAE Training process



[VAE Pretraining]

Action [0.1, 0.2, 0.3, ...] → “code #15”

VQ loss

“code #15” → Action Reconstruction [0.09, 0.21,..]

L1 loss

Step2

Map VQ Codes to LLM Token (manual mapping)

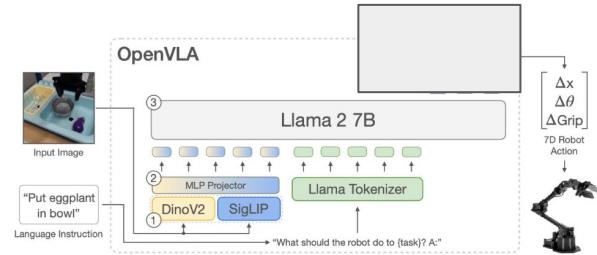
Level 1 code [0-255]
→ LLM Tokens [31040-31295]

Level 2 code [0-255]
→ LLM Tokens [31296-31551]

..

Step3

Fine-tune OpenVLA (frozen decoder) to predict action tokens



[OpenVLA]

Generate LLM Tokens before action de-tokenizers: [31045, 31302, ..]

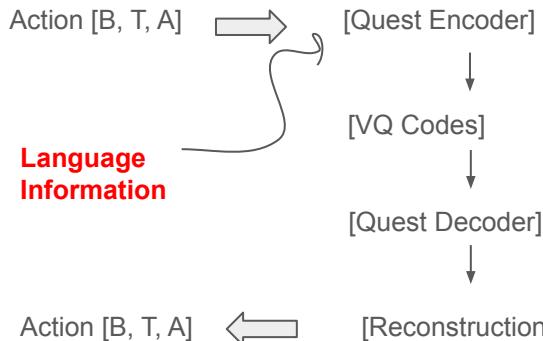
VQVAQActionTokenizer.
decode()

[Action]

Ours

Step1

Quest + Conditioned Training process



Step2

Map VQ Codes to LLM Token (manual mapping)

Encode: Action \rightarrow VQ Codes \rightarrow LLM Tokens

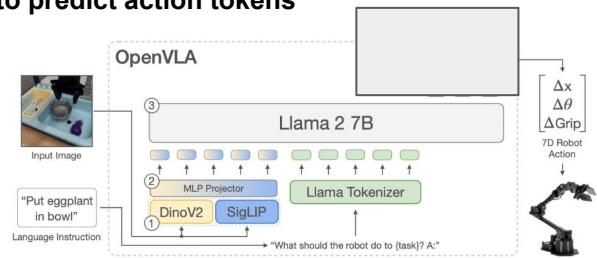
- Quest.encode \rightarrow Codes
- Code Mapping
- LLM Tokens

Decode: Tokens \rightarrow VQ Codes \rightarrow Action

- Token - 32032 \rightarrow Code
- Code to action

Step3

Fine-tune OpenVLA (frozen quest decoder) to predict action tokens



[OpenVLA]

Generate LLM Tokens before action
de-tokenizers: [31045, 31302, ...]

ours.decode()



[Action]



Quantitative Metric

Current Status

TODO

1. Integration Quest Decoder into VLA
2. Fine-tuning VLA with Conditioned Quest Decoder
3. Metric Evaluation
4. Simulation & Real-robot Task

2. Metric

Motivation

Beyond basic reconstruct loss and qualitative analysis e.g., skill was well compressed, let's create "**Quantitative Criteria**" for skill abstraction quality to evaluate environments (models, data, etc)

High Level Understanding

Prior work shows that **tokens can reconstruct and control**, **but they don't tell us**

- (i) **Semantic alignment**, “does a token consistently connect one motion concept?”
- (ii) **Composability**, “is that meaning invariant across episodes/tasks?”
- (iii) **compression**, “does this token is efficiently compressed?”

So, we are adding **alignment** to verify semantic meaning and **composability** to find episode-invariance, **BPA** for compression, and so on.

These metrics demonstrate that our learned tokens are **reusable, compact, and somehow interpretable** exactly what a skill library should be.

Metric - Few Details

Prior work: what they measure (Quest, Vq-BET)

- **Reconstruction / NLL**: how well actions decode from tokens.
- **Perplexity**: effective number of codes used.
- **Task return / success rate**: downstream control performance.
- **Occasional visuals**: nearest-neighbor rollouts or token retrievals.

But these miss essentials for *skills*:

1. **Semantic alignment** — *does a token consistently mean one motion concept?*
2. **Composability** — *is that meaning invariant across episodes/tasks?*

Why this is important

- **Interpretability**: tokens are *meaningful skills, not arbitrary chunks* (High alignment)
- **Transfer & reuse**: the same token works *when & where you need it* (High composability)
- **Systems benefit**: *cheaper planning, storage, and communication without control degradation* (Low BPA + good recon)

Our evaluation: what we add

A. Compression / efficiency

- **Bits-per-action (BPA)**: fewer bits at equal/better recon = better rate–distortion.
- **Compression ratio / token length**, plus **codebook utilization** sanity checks.

B. Segment-level “alignment” (InfoMEC family)

- **infomec_align_modularity**: each code specializes to one motion type.
- **infomec_align_compactness**: each motion type uses few codes (fewer synonyms).
- **infomec_align_explicitness**: linear probe predicts motion from codes.

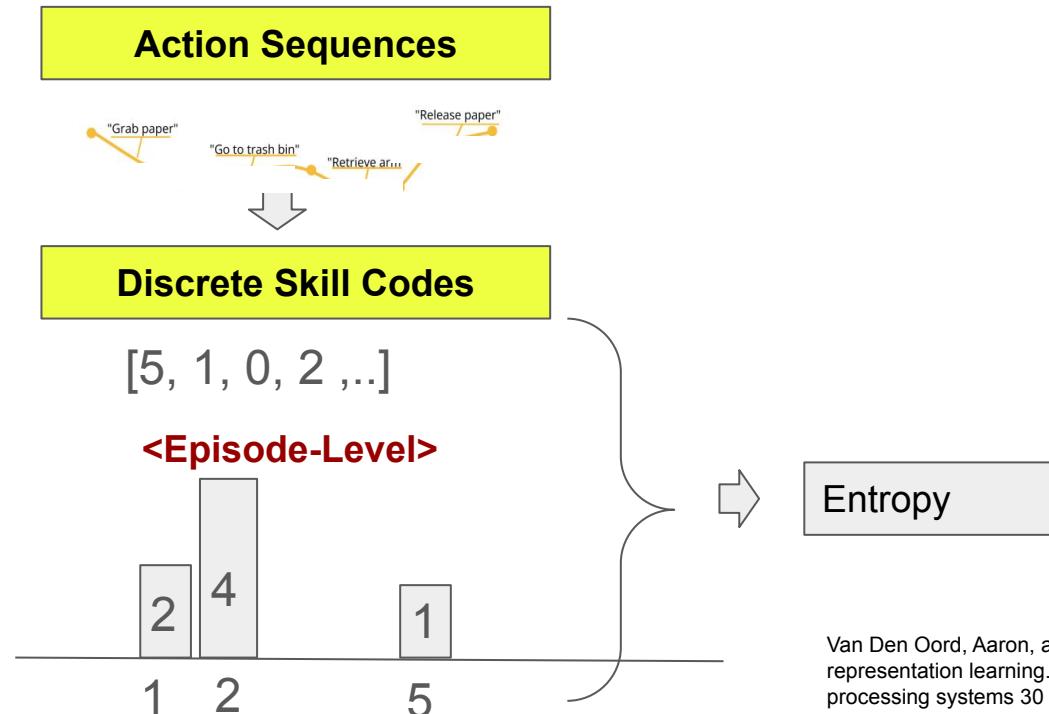
C. Sequential (Ongoing)

- **discrete first-order Markov model**

Bits-per-action (compactness)

How many bits we need to encode this sequence? Want to know skill primitives

- Episode-level evaluation, If not, it might correctly measure each metric since it will obtain value from randomly mixing skill information.



Van Den Oord, Aaron, and Oriol Vinyals. "Neural discrete representation learning." Advances in neural information processing systems 30 (2017).

Bits-per-action (compression)

Let say,

If use all 32 codes equally

- indices = [0,1,2,...,31,0,1,2,...,31, ...]
- unique_indices = [0,1,2,...,31]
- code_probs = [1/32, 1/32, , ...] # Uniform distribution
- entropy = $\log_2(32) = 5.0$ bits per action

If use 16 codes

- indices = [0,1,2,...,15,0,1,2,...,15, ...]
- unique_indices = [0,1,2,...,15]
- code_probs = [1/16, 1/16, ..., 0, 0, ...] # Half unused
- entropy = $\log_2(16) = 4.0$ bits per action

Episode-level

Avoiding randomly mixing skill information to complete one episode

Episode 1: VQ codes [3,1,7,2] → entropy = 2.0 bits (4 different codes)

Episode 2: VQ codes [3,1,7,2] → entropy = 2.0 bits (4 different codes)

Episode 3: VQ codes [5,8,4,9] → entropy = 2.0 bits (4 different codes)

Episode 4: VQ codes [1,1,1,1] → entropy = 0.0 bits (1 repeated code)

..
Average: $(2.0 + 2.0 + 2.0 + 0.0 + \dots) / \text{num_episodes} = 1.456$

To get entropy of codebook,

1. Count how often each code is used

```
code_counts = torch.bincount(predicted_skill_id, minlength=32)  
# code_counts = [85, 12, 45, 67, 23, 8, ...] (32 values)
```

2. Compute probability distribution of code counts

```
code_probs = code_counts / code_counts.sum()
```

3. Get Entropy (bits per action)

```
entropy = -torch.sum(code_probs * torch.log2(code_probs + 1e-10))
```

Assume that 4 bit is proper points for compression,
~16 codes used (medium entropy = 4.0 bits)

If entropy ≈ 4.0 , we're using ~lower codes efficiently

If entropy > 5.0 , we're using too many codes (inefficient)

If entropy < 3.0 , we're using too few codes (under-expressive)

Bits-per-action (compression)

```
186     def compute_bits_per_action(
187         self,
188         indices: torch.Tensor,
189         codebook_size: int
190     ) -> float:
191         """
192             Compute bits-per-action compression metric.
193
194         Args:
195             indices: (B, T) codebook indices
196             codebook_size: size of codebook K
197
198         Returns:
199             Average bits per action (lower = better compression)
200         """
201         if hasattr(indices, 'cpu'):
202             indices = indices.cpu().numpy()
203
204         indices_flat = indices.reshape(-1)
205
206         # Compute empirical distribution over codebook indices
207         unique_indices, counts = np.unique(indices_flat, return_counts=True)
208         probs = counts / len(indices_flat)
209
210         # Compute entropy in bits
211         entropy_bits = -np.sum(probs * np.log2(probs + 1e-12))
212
213         return float(entropy_bits)
```

VAE output

trajectory input -> VAE encoder -> ResidualVQ ->
Codes (=indices)

Count how often each code is used

Entropy

InfoMEC

What is InfoMEC?

InfoMEC measures how well learned representations (like VQ codes) capture meaningful structure in data.

- Do the codes we learned actually represent distinct, useful concepts?

The Three Components:

1. InfoMEC-M (Modularity): "Does each code specialize in one thing?"

- Good: Code A = "grasping", Code B = "moving"
- Bad: Code A = "grasping + moving + rotating"

2. InfoMEC-E (Explicitness): "Can we predict the task from the codes?"

- Good: Seeing codes [A,B,C] → I know it's "pick and place"
- Bad: Codes don't tell me what task is happening

3. InfoMEC-C (Compactness): "Does each concept use few codes?"

- Good: "Grasping" always uses codes [A,B]
- Bad: "Grasping" scattered across codes [A,B,C,D,E,F]

InfoMEC - Original paper

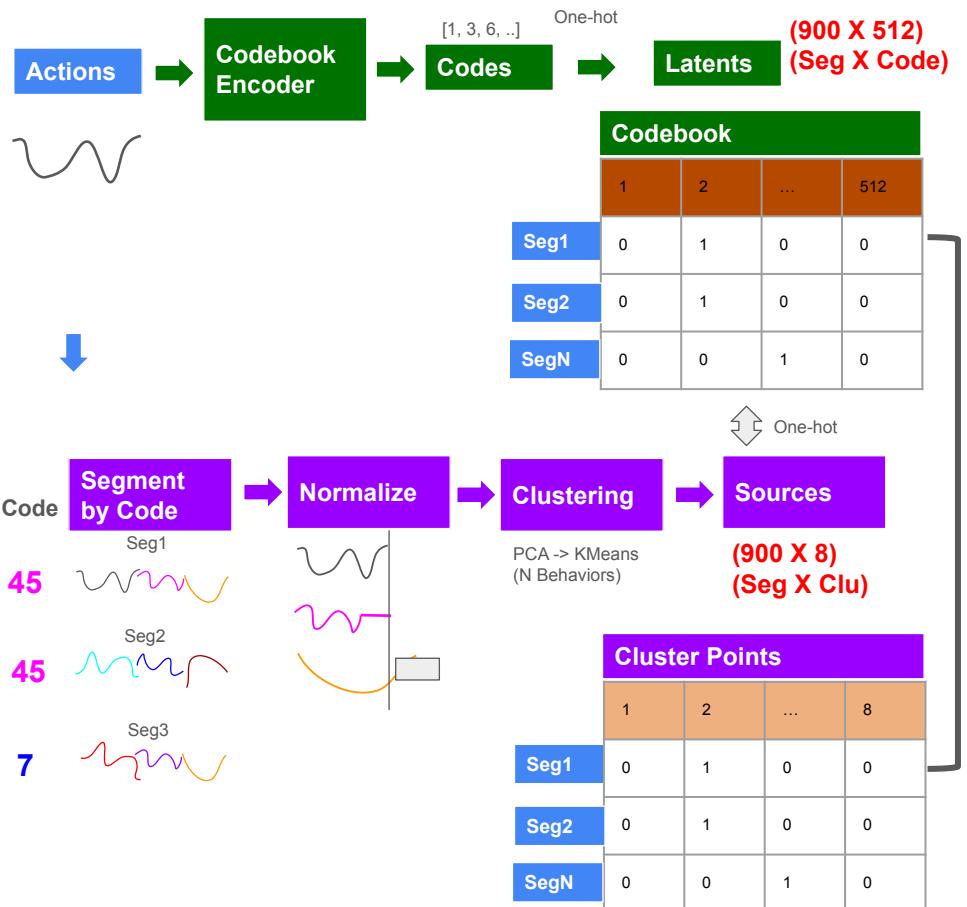
- The **sources** s_i are **given by the datasets** (not computed by binning raw pixels). Examples: “object color (10 values)”, “robot-x (40 bins)”, “camera height (4)”, etc. Those datasets (Shapes3D, MPI3D, Falcor3D, Isaac3D) already **provide discrete factor labels** **Note: the datasets already define discrete factors.**
- The **latents** z_j are **discrete** because the encoder output is **quantized per dimension** to a small scalar **codebook** (e.g., 10 values).
- InfoMEC is then **computed between ground-truth discrete factors and the model's discrete latents.** (If any variable were continuous, they'd use a KSG MI estimator; but most factors there are discrete.)

One Issues

- In our case, we don't have any discrete labels of sources, so we need to use unsupervised way to give some labels.

InfoMEC - Alignment

how well cluster patterns can be represented by vq codes.



Co-matrix (NMI) (8 X 512)

Codebook		
C1	167	...
C2	7	135
C3	0	760
C4	1	...
C5	2	2
C6	10	1
C7	58	530
C8	156	100

NMI

Codebook		
C1	0.85	0.02
C2	0.03	0.89
C3	0.01	0.01
C4	0.01	..
C5	0.01	0.02
C6	0.03	0.01
C7	0.2	0.8
C8	0.5	0.2

InfoMEC-M

- Column-wise (Codebook)
1. col_ratios = **col_max** / col_sums
 2. col_mean = np.mean(col_ratios)
 3. InfoMEC-M= (mean - 1/clusters) / (1 - 1/clusters)

See whether each VQ code specializes in one behavioral pattern

InfoMEC-C

- Row-wise (Cluster)
1. row_ratios = row_max / **row_sums**
 2. row_mean = np.mean(row_ratios)
 3. InfoMEC-C formula: (mean - 1/codebook) / (1 - 1/codebook)

See whether each behavioral pattern uses few VQ codes

InfoMEC-E

Logistic Regression (X -> Y)

```
infoe = compute_infoe(sources, latents,
discrete_latents=True)
# X = latents (987, 512)
# y = behavioral_clusters (987,)
```

See whether VQ codes are predictive of behavioral patterns

InfoMEC - Alignment (details)

Input Data:

```
Trajectories = (32, 256, 7)      # 32 episodes, 256 timesteps, 7D actions  
indices.shape = (32, 256)        # 32 episodes, 256 VQ codes each
```

VQ code (256 codes) - 1st episode

```
indices[0] = [45, 45, 45, 45, 127, 127, 127, 234, 234, 89, ..]
```

Step 1: Extract segments from ALL episodes together

```
# Collect segments across all episodes, ~800-1200 total segments collected
```

Step 2: Normalize segment lengths: Truncate or Right-Zero Padding

```
# 987 segments × 14D (target_length)
```

Step 3: PCA and clustering

```
# StandardScaler + PCA (987, 50) -> KMeans clustering, capped at 8 clusters  
behavioral_clusters.shape = (987,) # [0,1,2,3,4,5,6,7,0,1,2,...]
```

Step 4: Create matrices for InfoMEC

```
# Sources: behavioral clusters one-hot encoded: # 987 segments × 8 behavioral patterns  
# Latents: VQ codes one-hot encoded) # 987 segments × codebook_size = 512
```

Step 5: Compute InfoMEC Metrics

```
A. Compute NMI Matrix
```

```
result = compute_infomec_original()
```

```
sources=sources,      # (987, 8) behavioral clusters  
latents=latents,     # (987, 512) VQ codes
```

```
)
```

Step 6. Matrix Build

```
for segment_idx in range(987):  
    pattern_id = np.argmax(sources[i]) # Which pattern? (0-7)  
    vq_code_id = np.argmax(latents[i]) # Which VQ code?
```

```
# Result:
```

```
co_occurrence_matrix = [  
    [167, 5, 2, 8, 18, ...], # code45=167x, code89=5x,  
    [ 7, 134, 3, 4, 12, ...], # code45=7x, code89=134x,  
    [ 3, 2, 142, 6, 7, ...], # code127=142x  
    # ... 5 more patterns  
]
```

Step 7: Calculate probabilities and MI

```
P_joint = co_occurrence_matrix / total_segments  
P_pat = np.sum(P_joint, axis=1) # Marginal: P(pattern)  
P_vq_c = np.sum(P_joint, axis=0) # Marginal: P(vq_code)
```

```
P_pat[0] = 0.20 # 20% of seg are pattern 0  
P_vq_c[0] = 0.18 # 18% of seg use VQ code 45  
P_joint[0,0] = 0.17 # 17% of seg both 0 AND code 45
```

```
for i in range(8):    # For each behavioral pattern  
    for j in range(): # For each VQ code  
        MI[i,j] = P_joint[i,j] * np.log2(P_joint[i,j] / (P_pattern[i]  
* P_vq_code[j]))
```

Step 8: Normalize by entropy

```
# Normalized MI: NMI(X,Y) = MI(X,Y) / H(X)  
# H(X) = entropy of behavioral patterns  
H_pattern = -np.sum(P_pat * np.log2(P_pat + 1e-10))
```

InfoMEC - Alignment (details)

A. Example NMI matrix:

```
nmi = np.array([
    [0.85, 0.02, 0.01, 0.03, 0.09, ...], # Behavioral pattern 0 → mainly VQ code 45
    [0.03, 0.89, 0.01, 0.02, 0.05, ...], # Behavioral pattern 1 → mainly VQ code 89
    [0.02, 0.01, 0.91, 0.03, 0.03, ...], # Behavioral pattern 2 → mainly VQ code 127
    [0.01, 0.03, 0.02, 0.88, 0.06, ...], # Behavioral pattern 3 → mainly VQ code 234
    [0.04, 0.02, 0.01, 0.04, 0.89, ...], # Behavioral pattern 4 → mainly VQ code 312
    [0.02, 0.01, 0.02, 0.01, 0.02, ...], # Behavioral pattern 5 → mixed codes
    [0.01, 0.02, 0.01, 0.02, 0.01, ...], # Behavioral pattern 6 → mixed codes
    [0.02, 0.01, 0.01, 0.01, 0.01, ...] # Behavioral pattern 7 → mixed codes
])
```

B. InfoMEC-M (Modularity) - Column-wise

```
# For each VQ code (column), check if it's modular
col_max = np.max(nmi, axis=0) # [0.85, 0.89, 0.91, 0.88, 0.89, ...] # 60 values
col_sums = np.sum(nmi, axis=0) # [1.00, 1.00, 1.00, 1.00, 1.00, ...] # normalized
col_ratios = col_max / col_sums # [0.85, 0.89, 0.91, 0.88, 0.89, ...]
col_mean = np.mean(col_ratios) # 0.87

# InfoMEC-M formula: (mean - 1/num_sources) / (1 - 1/num_sources)
num_sources = 8
infom = (0.87 - 1/8) / (1 - 1/8) = (0.87 - 0.125) / 0.875 = 0.85
# High modularity: each VQ code specializes in one behavioral pattern
```

C. InfoMEC-C (Compactness) - Row-wise

```
# For each behavioral pattern (row), check if it's compact
row_max = np.max(nmi, axis=1) # [0.85, 0.89, 0.91, 0.88, 0.89, ...]
row_sums = np.sum(nmi, axis=1) # [1.00, 1.00, 1.00, 1.00, 1.00, ...]
row_ratios = row_max / row_sums # [0.85, 0.89, 0.91, 0.88, 0.89, ...]
row_mean = np.mean(row_ratios) # 0.82

# InfoMEC-C formula: (mean - 1/num_active_latents) / (1 - 1/num_active_latents)
num_active_latents = 60
infoc = (0.82 - 1/60) / (1 - 1/60) = (0.82 - 0.017) / 0.983 = 0.82
# High compactness: each behavioral pattern uses few VQ codes
```

D. InfoMEC-E (Explicitness) - Logistic Regression

```
# Can we predict behavioral patterns from VQ codes?
infoe = compute_infoe(sources, latents, discrete_latents=True)

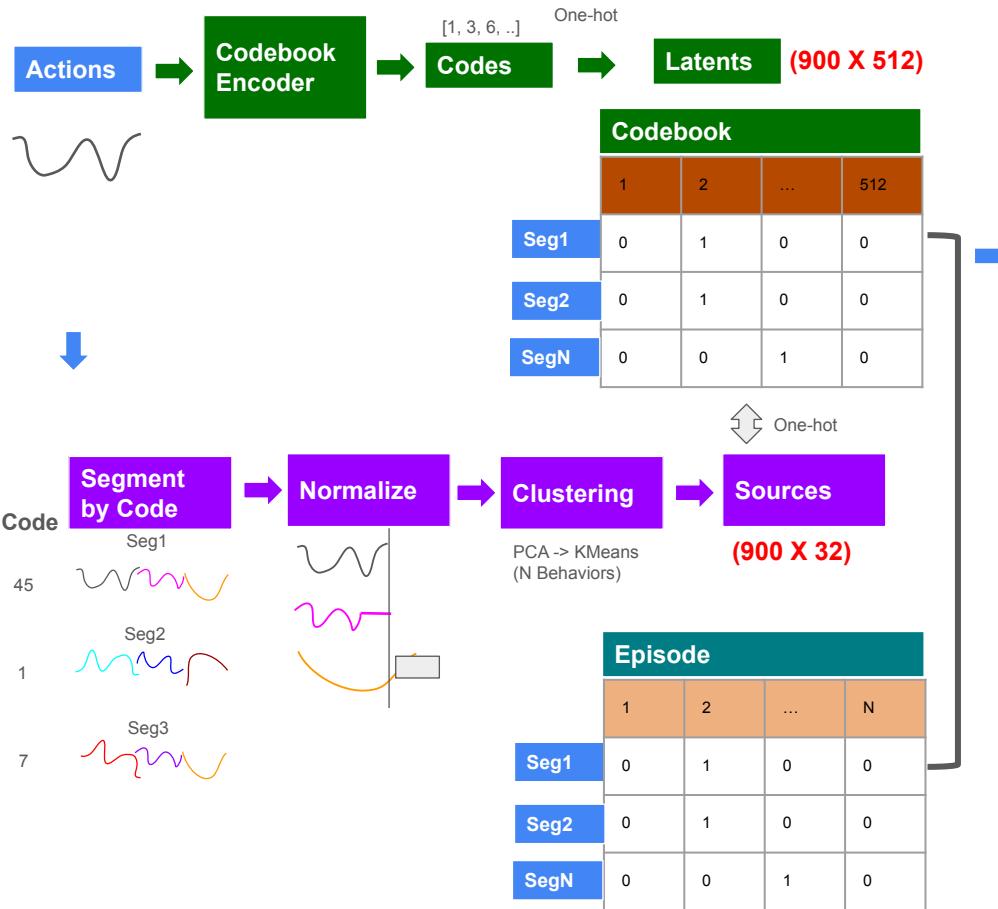
# X = latents (987, 512) - VQ code features
# y = behavioral_clusters (987,) - target labels [0,1,2,3,4,5,6,7,0,1,...]

predictive_loss = logistic_regression(latents, behavioral_clusters)
null_loss = logistic_regression(np.zeros_like(latents), behavioral_clusters)

infoe = (null_loss - predictive_loss) / null_loss = (2.08 - 0.25) / 2.08 = 0.88
# High explicitness: VQ codes are predictive of behavioral patterns
```

InfoMEC-C - Composability

how well skill patterns can be reusable across episodes



Co-matrix (NMI)			
Codebook			
E1	167	...	1
E2	7	...	135
E3	0	...	760
E4	1
E5	2	...	2
E6	10	...	1
E7	58	...	530
N	156	...	100

InfoMEC-C

Row-wise
(Cluster)

1. `row_ratios = row_max / row_sums`
2. `row_mean = np.mean(row_ratios)`
3. InfoMEC-C formula: $(\text{mean} - 1/\text{codebook}) / (1 - 1/\text{codebook})$

See whether each behavioral pattern uses few VQ codes

Codebook			
E1	0.85	0.02	..
E2	0.03	0.89	..
E3	0.01	0.01	...
E4	0.01
E5	0.01	..	0.02
E6	0.03	..	0.01
E7	0.2	..	0.8
N	0.5	..	0.2

Epi 1: uses codes 45,89,127,234,312

Epi 2: uses codes 45,89,127,234,312

Epi 3: uses codes 45,89,127,234,312

Epi 4: uses codes 45,89,127,234,312

... 28 more Epi

Epi 32: balanced usage

InfoMEC-C - Composability (details)

Step 1. Matrix Build

```
for segment_idx in range(987):
    pattern_id = np.argmax(sources[i]) # Which pattern? (0-7)
    vq_code_id = np.argmax(latents[i]) # Which VQ code?

co_occurrence_matrix = [
    [8, 7, 6, 8, 9, 0, 0, 0, ...], # Epi 0: uses codes 45,89,127,234,312
    [6, 9, 8, 7, 8, 0, 0, 0, ...], # Epi 1: uses codes 45,89,127,234,312
    [7, 8, 9, 6, 7, 0, 0, 0, ...], # Epi 2: uses codes 45,89,127,234,312
    [9, 6, 7, 9, 8, 0, 0, 0, ...], # Epi 3: uses codes 45,89,127,234,312
    # ... 28 more Epi
    [8, 8, 8, 8, 8, 0, 0, 0, ...], # Epi 31: balanced usage
]
```

Step 2. Calculate probabilities and MI

```
P_episode = [0.031, 0.031, 0.031, ...] # Each episode ~3.1%
P_vq_code = [0.18, 0.19, 0.17, 0.16, 0.18, ...] # VQ code distribution
P_joint[0,0] = 8/987 = 0.008 # P(episode=0 AND vq_code=45)

for i in range(32):      # For each episode
    for j in range(N):   # For each VQ code
        if P_joint[i,j] > 0: # Avoid log(0)
            MI[i,j] = P_joint[i,j] * np.log2(P_joint[i,j] / (P_episode[i] *
P_vq_code[j]))
```

Step 3. Normalize by entropy

```
H_episode = -np.sum(P_episode * np.log2(P_episode + 1e-10))
```

Step 4: Compute InfoMEC-C only (composability metric)

```
# NMI between 32 episodes and active VQ codes
nmi.shape = (32, 60)

# Example episode-level NMI:
episode_nmi = np.array([
    [0.20, 0.18, 0.22, 0.15, 0.25, ...], # Epi 0: codes somewhat equally
    [0.25, 0.20, 0.20, 0.15, ...], # Epi 1: codes somewhat equally
    [0.15, 0.30, 0.25, 0.15, 0.15, ...], # Epi 2: for some codes
    # ... 29 more episodes
])
```

InfoMEC-C for episodes (lower = better composability)

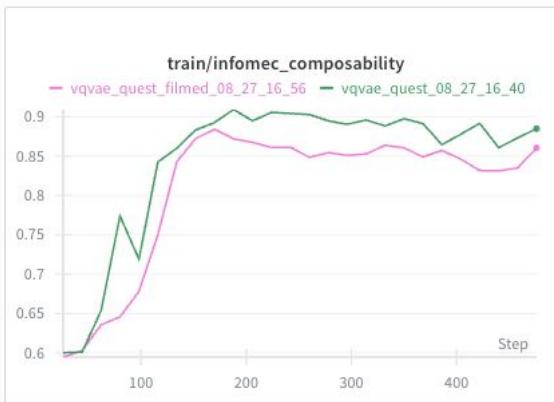
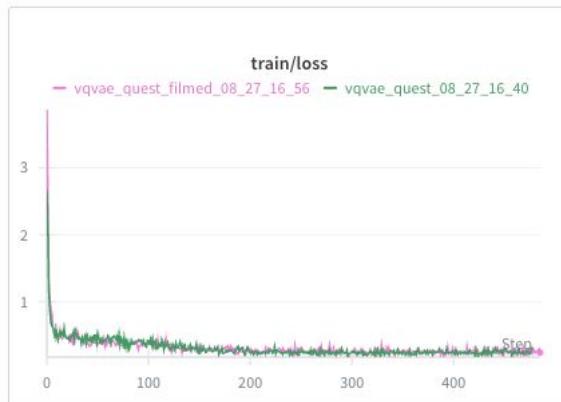
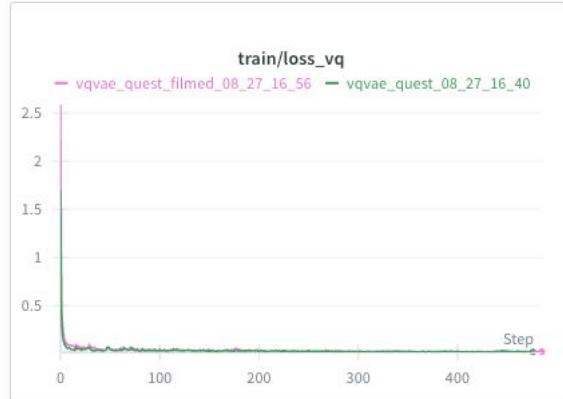
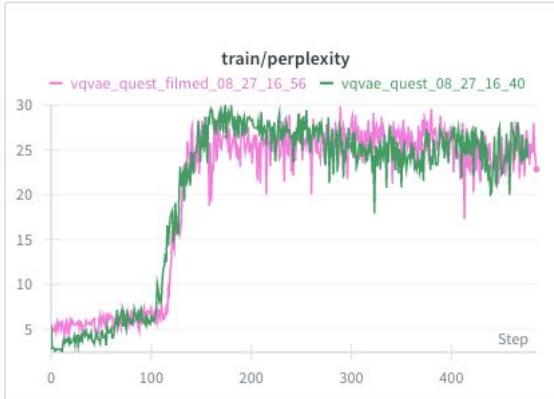
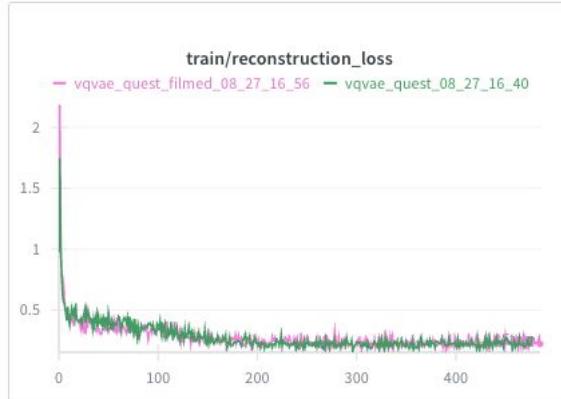
```
row_max = np.max(episode_nmi, axis=1) # Max NMI per episode
row_sums = np.sum(episode_nmi, axis=1) # Sum per episode
row_ratios = row_max / row_sums # Ratio per episode
row_mean = np.mean(row_ratios) # Average ratio
```

```
num_active_latents = 60
episode_infoc = (row_mean - 1/60) / (1 - 1/60) # e.g., 0.35
```

Transform to composability score

```
composability_score = max(0.0, 1.0 - episode_infoc) # e.g., 0.65
```

Results - 1



Results - 2



Key Metrics & Observations

Alignment dimensions

- **Modularity** → Each code specializes in specific movement patterns
- **Explicitness** → Behavior can be predicted from VQ codes
- **Compactness** → Patterns map to focused VQ codes

Results

- **Loss**: 0.245 vs 0.280 (**-12%**) → better training
- **Perplexity**: 22.85 vs 24.48 (**-7%**) → effective codebook usage, fewer codes used
- **Modularity (alignment)**: 0.362 vs 0.321 (**+13%**) → VQ codes specialize in specific movement patterns
- **Explicitness (alignment)**: 0.327 vs 0.306 (**+7%**) → codes more predictive
- **Compactness (alignment)**: 0.196 vs 0.168 (**+17%**) → fewer codes used?
- **BPA (bits per action)**: 3.571 → 3.308 (**-7.4%**) → model needs ~0.26 fewer bits per action → better representations

Note:

- alignment improves (codes become more specialized/meaningful)

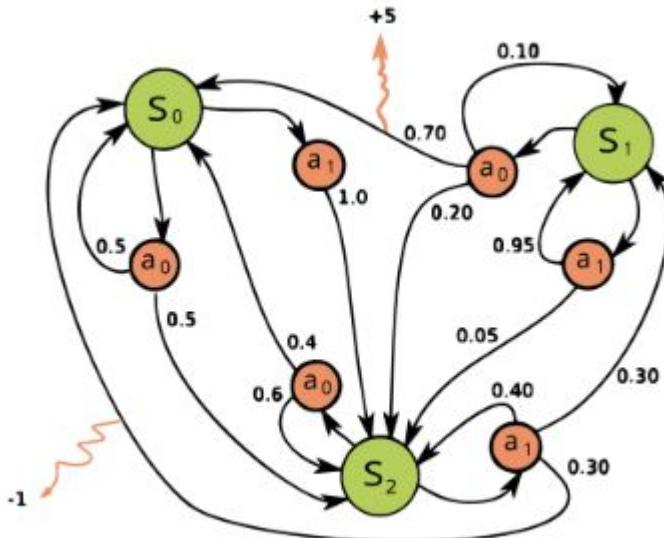
Markov Decision Process (MDPs)

Markov Decision Processes (MDPs)

MDPs are a mathematical model for sequential decision-making in a fully observable, stochastic, discrete-time environment with a Markovian transition model.

$$M = \langle S, A, T, r \rangle$$

$$A = \{\leftarrow, \uparrow, \rightarrow\}$$



MDP Metric Objective

1) Episode-level (meso): **Sequentiality / predictability**

What it answers: *How predictable is the next code from the last K codes within an episode?*

Why this level: To understand strong sequential structure & sequence property.

2) Task-level (macro): **Composability / invariance**

What it answers: *Do different episodes/tasks use the same tokens and chain them the same way?*

Why this level: Reuse/transfer is about **consistency across tasks/episodes**.

discrete first-order Markov model

$P(\text{next_code} \mid \text{prev_code})$ - First-order Markov

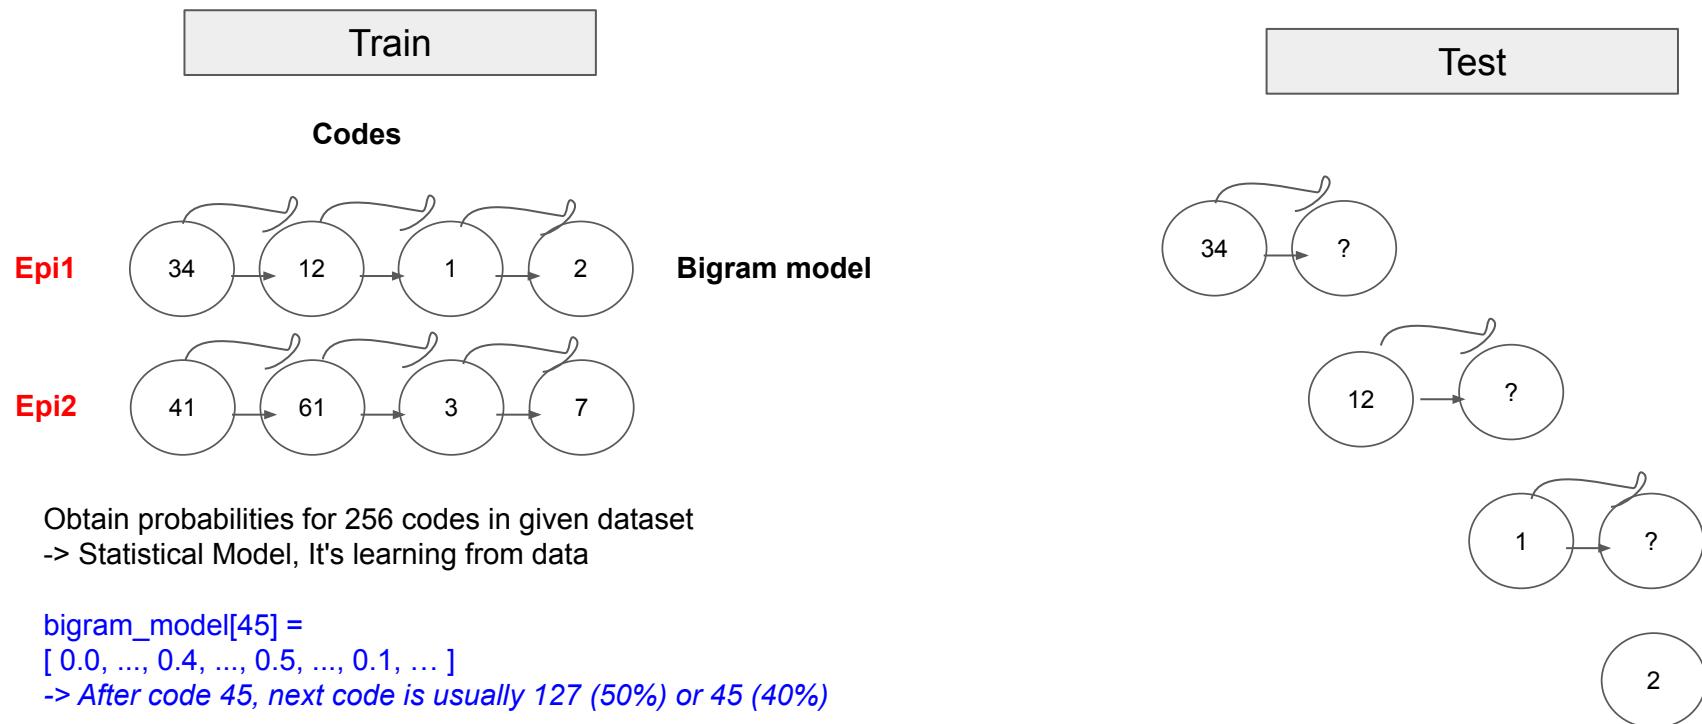
- **Markov assumption:** The next state depends **only** on the current state, not the full history.
- **No control** (passive observation) & **No Action Space** [Action = state (predict next code, which becomes next state)]
- **Deterministic policy:** Always pick argmax

MDP Component	Our Implementation
State s_t	prev (last 1 token, K=1)
Action a_t	argmax(model[prev]) (<i>pick top-1 next code</i>)
Transition $P(s_{t+1} s_t)$	bigram_model[prev, next], # <i>Learning $P(z_{t+1} z_t)$ from data</i>
Reward $r(s,a,s')$	1 if pred==true else 0 (correctness)
Discount γ	1 (finite episodes, average per step)
Discrete	256 codes (finite state space)

Sequentiality

VQ model captures temporal structure in the behavior?

Learn bigram model $P(\text{next_code} \mid \text{prev_code}) \rightarrow$ Test: Can we predict next code?



1. Sequentiality

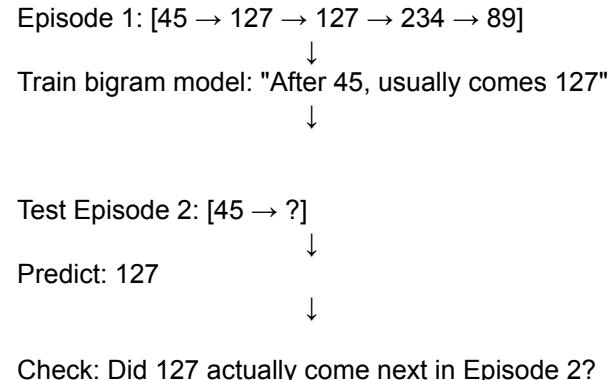
how well VQ model captures temporal structure in the behavior

What's this step: Train data → Learn $P(\text{next_code} | \text{prev_code})$ → Test: Can we predict next code?

How we get:

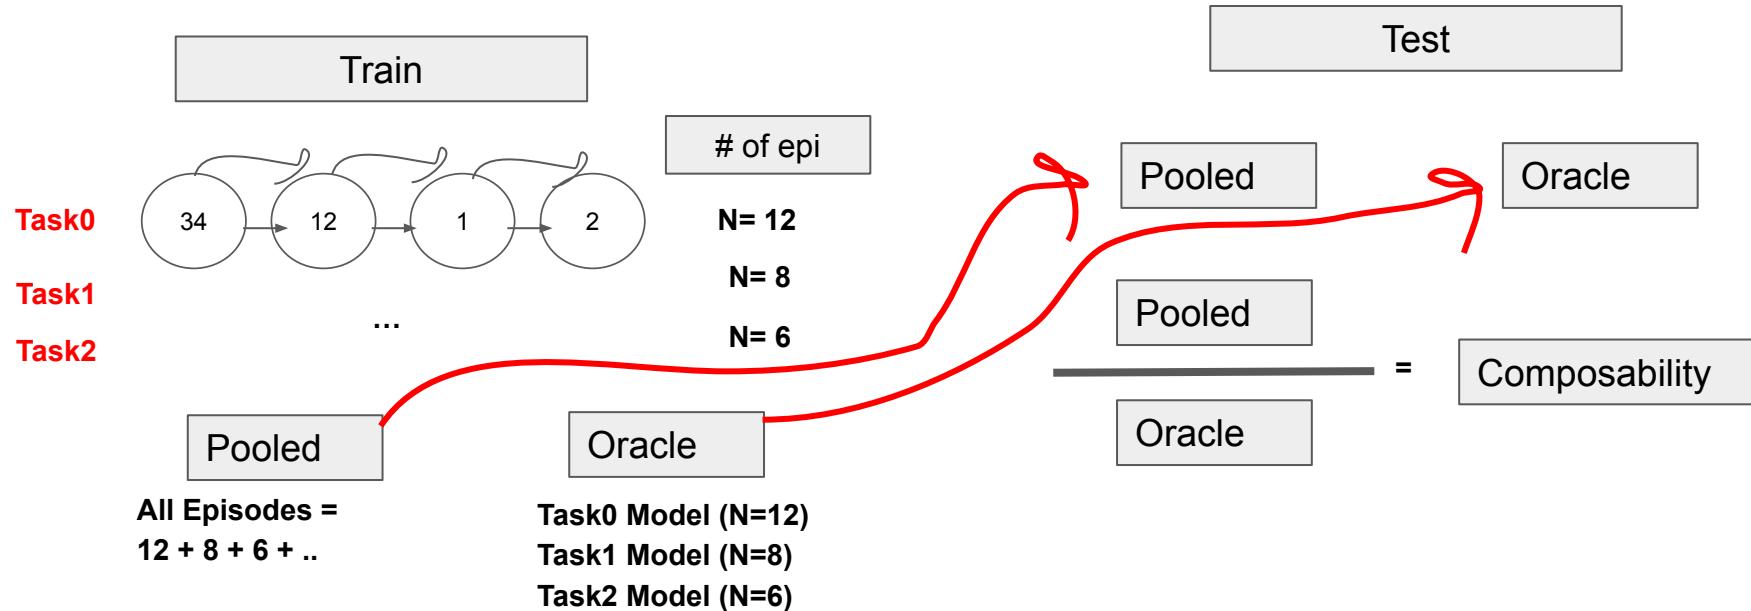
```
def evaluate_predictability(test_eps, model):
    for ep in test_eps:
        for prev, true_next in zip(ep[:-1], ep[1:]):
            pred_next = np.argmax(model[prev])
            if pred_next == true_next:
                correct += 1
```

- **It's a Statistical Model, It's learning from data**
 - Train on train episodes, test on held-out test episodes
- **Actual next word is the "ground truth"**
 - Given code 45 at time t, can you predict that code came next?
 - The **observed sequence** is the ground truth



Short summary - Composability

Do different tasks reuse the same VQ code patterns?



2. Composability

Do different tasks reuse the same VQ code patterns?

→ "How much do tasks share patterns?

What's this step: Compare: One general model vs Task-specific models

How we get: **Train pooled and oracle model**

Assumption

1. **High sharing** -> we can learn one model for multiple tasks (efficient)
2. **Low sharing** -> Each task needs its own specialized knowledge (harder to generalize)

So setup,

1. **Generalist performance:** How well does one general course work for everyone?
 - a. **Pooled:** Cross-task generalization -> Can one model handle everything?
2. **Specialist performance:** How well do specialized courses work?
 - a. **Oracle:** Best possible performance -> How well can we do if specialized?

Ratio: Pattern sharing -> "How much do tasks share patterns?"

If **generalist ≈ specialist** → General course is fine! (High composability)

If **generalist << specialist** → Need specialized courses! (Low composability)

Short summary - Composability

High Composability (pooled ≈ oracle)

Task 0: [45→127→234→89] # reach, grasp, lift, place
Task 1: [45→127→234→50] # reach, grasp, lift, place (different object)
Task 2: [45→127→180→89] # reach, grasp, different motion, place

All share: 45→127 (reach→grasp)

Pooled model learns:

After 45 → predict 127 (100% accurate across all tasks)
After 127 → predict 234 (60%) or 180 (40%) - pretty good

Oracle model learns (Task 0 only):

After 45 → predict 127 (100%)
After 127 → predict 234 (100% for Task 0 specifically)

Result: Pooled is 95% as good as Oracle

- Because tasks share most transitions
- The pooled model benefited from seeing more examples of 45→127

Low Composability (pooled << oracle)

Task 0: [10→20→30→40] # Unique code sequence
Task 1: [100→110→120→130] # Completely different codes
Task 2: [200→210→220→230] # No overlap at all

Pooled model learns:

After 10 → predict ??? (only seen in Task 0, confused with other tasks)
After 100 → predict ??? (only seen in Task 1)
Accuracy: 35% - just random guessing

Oracle model learns (Task 0 only):

After 10 → predict 20 (100% - specialized)
After 20 → predict 30 (100%)
Accuracy: 95%

Result: Pooled is only 37% as good as Oracle

- Because tasks don't share anything
- The pooled model is harmed by mixing incompatible patterns

Another Intuition

If patterns are shared:

$P_{\text{pooled}}(\text{next}|\text{prev})$
 $\approx P_{\text{task0}}(\text{next}|\text{prev})$
 $\approx P_{\text{task1}}(\text{next}|\text{prev})$

All tasks follow similar rules
→ Pooled model learns the common rule
→ Performs almost as well as specialists

If patterns are unique:

$P_{\text{pooled}}(\text{next}|\text{prev})$
= average of different rules $P_{\text{task0}}(\text{next}|\text{prev})$
 $\neq P_{\text{task1}}(\text{next}|\text{prev})$
 $\neq P_{\text{task2}}(\text{next}|\text{prev})$

Each task has different rules
→ Pooled model learns a confused average
→ Much worse than specialists

2. Composability

Do different tasks reuse the same VQ code patterns?

DATA SPLIT:

```
TRAIN (48 episodes)
  Task 0: 12 episodes
  Task 1: 10 episodes
  Task 2: 15 episodes
  Task 3: 11 episodes

TEST (32 episodes)
  Task 0: 8 episodes
  Task 1: 7 episodes
  Task 2: 9 episodes
  Task 3: 8 episodes
```

TRAINING PHASE:

```
Pooled Model (1 model)
  ← train ALL 48 episodes
```

```
Oracle Models (4 models)
  Model_0 ← train Task 0 (12 eps)
  Model_1 ← train Task 1 (10 eps)
  Model_2 ← train Task 2 (15 eps)
  Model_3 ← train Task 3 (11 eps)
```

TESTING PHASE:

```
Pooled: Test ALL 32 test episodes
  with same pooled model
  → pooled_acc = 0.82
```

```
Oracle: Test each task separately
  Task 0 test (8 eps) → Model_0
  → acc_0 = 0.87
  Task 1 test (7 eps) → Model_1
  → acc_1 = 0.92
  Task 2 test (9 eps) → Model_2
  → acc_2 = 0.85
  Task 3 test (8 eps) → Model_3
  → acc_3 = 0.90
  → oracle_acc = 0.885 (weighted)
```

COMPOSABILITY = $0.82 / 0.885 = 0.93$

Pooled (trained on ALL train episodes)

- pooled_acc =
evaluate_predictability(test_eps,
bigram_model)

*Oracle (trained only on episodes similar
to each test episode)*

Key Idea:

- If **pooled_acc ~ oracle_acc** →
episodes share patterns (high
composability)
- If **pooled_acc << oracle_acc** →
each group has unique patterns
(low composability)

HIGH Composability (Reusable Patterns)

Task 0 (Pick & Place):

[reach → grasp → lift → move → place]
↓ ↓ ↓ ↓ ↓
VQ: 45 → 127 → 234 → 89 → 50

Task 1 (Open Drawer):

[reach → grasp → pull → open → release]
↓ ↓ ↓ ↓ ↓
VQ: 45 → 127 → 180 → 200 → 75

Task 2 (Press Button):

[reach → align → press → retract]
↓ ↓ ↓ ↓
VQ: 45 → 90 → 150 → 30

LOW Composability (Unique Patterns)

Each task has completely different patterns:

Task 0:

VQ: 10 → 20 → 30 → 40 → 50

Task 1:

VQ: 100 → 110 → 120 → 130 → 140

Task 2:

VQ: 200 → 210 → 220 → 230 → 240

Results

Goal: Visualize how VQ-VAE skill codes transition from one to another over time.

- Nodes = Skill codes (e.g., "grasp", "move", "release")
- Edges = Transition probabilities (e.g., "grasp" -> "move" with 80% probability)
- The graph shows which skills naturally follow each other in robot behavior
- Training Episodes -> VQ Codes -> Transition Matrix -> Graph Visualization
- Example:
- Time: t=0 t=1 t=2 t=3 t=4
- Codes: [5] → [5] → [12] → [3] → [5]

This creates transitions:

- Code 5 → Code 5 (self-loop), Code 5 → Code 12, Code 12 → Code 3, Code 3 → Code 5

It's computed by counting how often one skill code transitions to another in one episode, normalizing those counts to probabilities, then visualizing the most frequent ones.

e.g.: [5 → 5 → 12 → 3 → 5], estimate P(5→5)=0.75, P(5→12)=0.25, build a weighted graph (nodes = codes, edges = transition probabilities), and draw thicker edges for higher probabilities.

[Details]

Step 1: After Build Bigram Transition Model,

Result: Transition matrix where $\text{probs}[5, 12] = 0.23$ means "after code 5, code 12 appears 23% of the time"

Step 2: Select Top Codes to Visualize

[5, 12, 3, 27, 18, ...] - the 15 most-used codes

Step 3: Build Graph Structure (Create DAG)

Example graph:

- Nodes: {3, 5, 12, 18, 27}

Edges:

5 → 5 (weight=0.65)

5 → 12 (weight=0.23)

12 → 3 (weight=0.51)

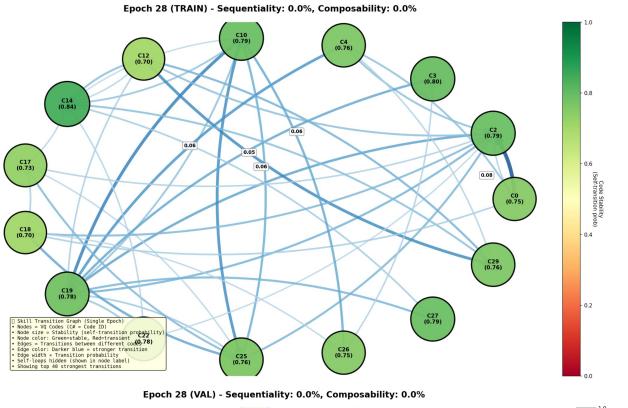
3 → 5 (weight=0.34)

Count transitions in training data → transition matrix -> Select top codes to avoid clutter -> Build graph with codes as nodes, transitions as edges

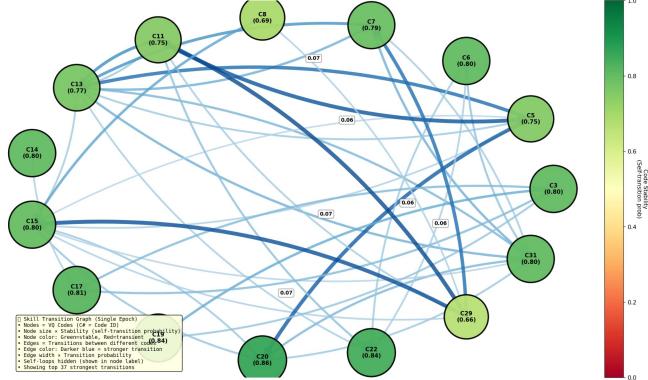
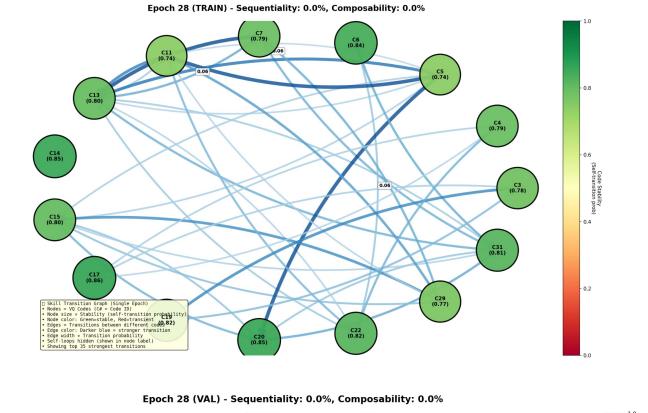
Result - Sequentiality

this is episode level epoch results (the end of one episode wouldn't be a point to connect with next episode initial point)
And, randomly selected epoch and task types amongst all tasks for composability

Baseline

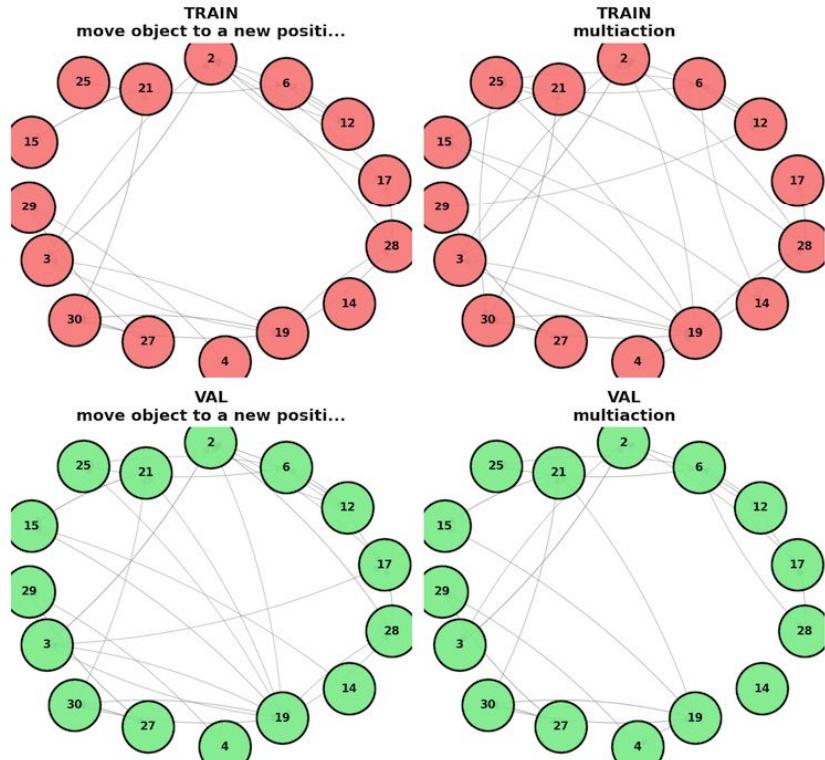


Film

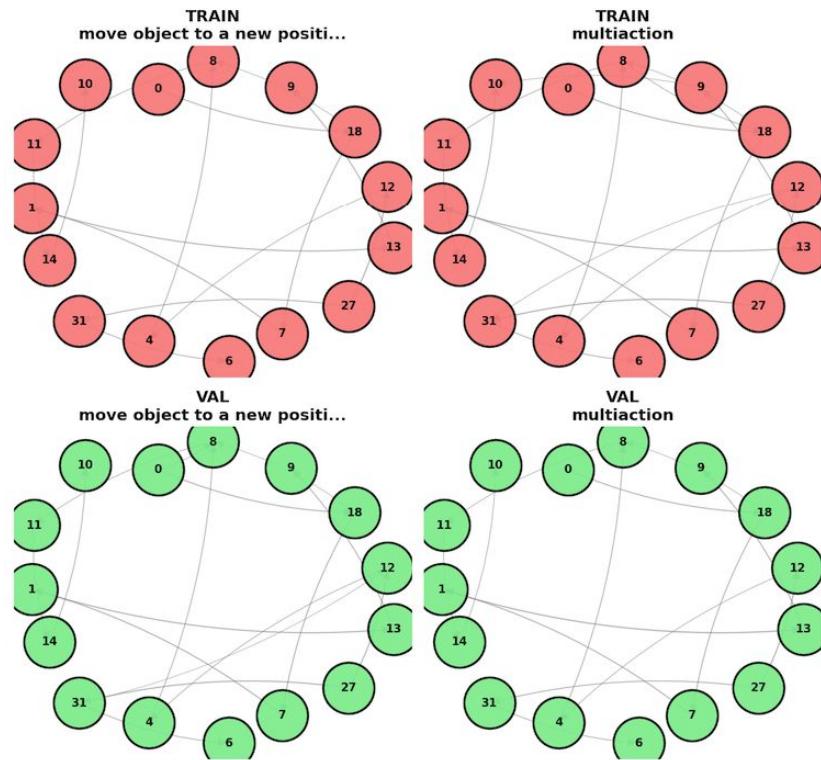


Result - Composability

Baseline



Film



Discussion

collapsed structure in both sequentiality and composability graphs suggests that transitions have become **more focused and deterministic**, meaning fewer random transitions and stronger skill-to-skill predictability.

- Some skill codes dominate transitions, implying the model has learned **stable, reusable primitives**.

- Stronger, thicker edges (high transition probabilities) show **clear temporal structure** — one skill consistently follows another

show **strong learned structure** (good), but **low cross-task transfer** (potentially limiting)

Sequentiality: Film > Baseline

- Film creates more deterministic, structured sequences
- Good train-val consistency = generalizes well

Related Questions

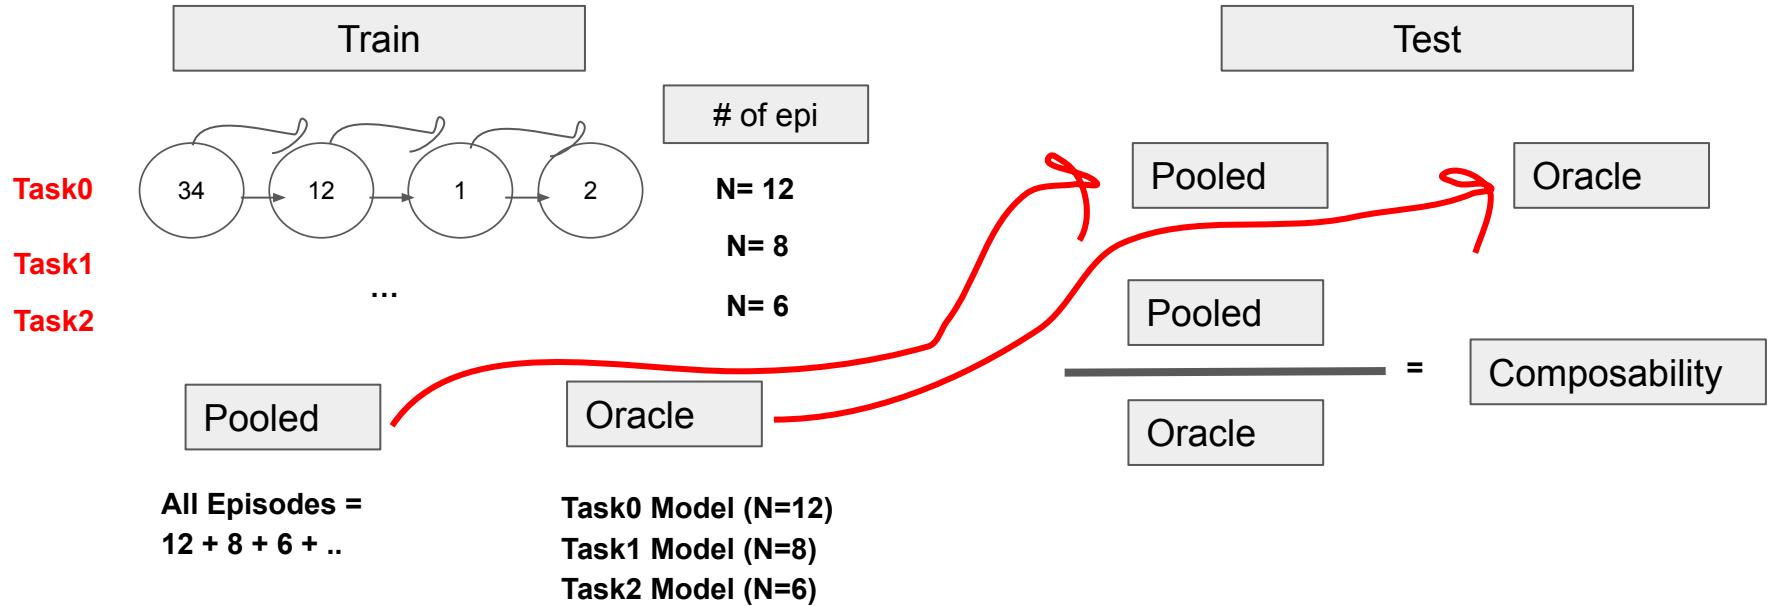
- Does higher sequentiality improve robot performance?
- Can composability be increased without reducing sequentiality?
- How does diversity vs. determinism trade off—does collapse imply overfitting?
- Quantify with entropy/sparsity of transitions and track edge dominance over epochs to monitor convergence.

3. Evaluation

- 1) Quest (baseline)**
- 2) Quest + Conditioned (our)**
- 3) VQ-VLA**
- 4) π model**

4. Experiments - Simulation (Libero) & Robot

Short summary - Composability



Pooled (trained on ALL train episodes)

- pooled_acc = evaluate_predictability(test_eps, bigram_model)

Oracle (trained only on episodes similar to each test episode)

- If $\text{pooled_acc} \approx \text{oracle_acc}$ → episodes share patterns (high composability)

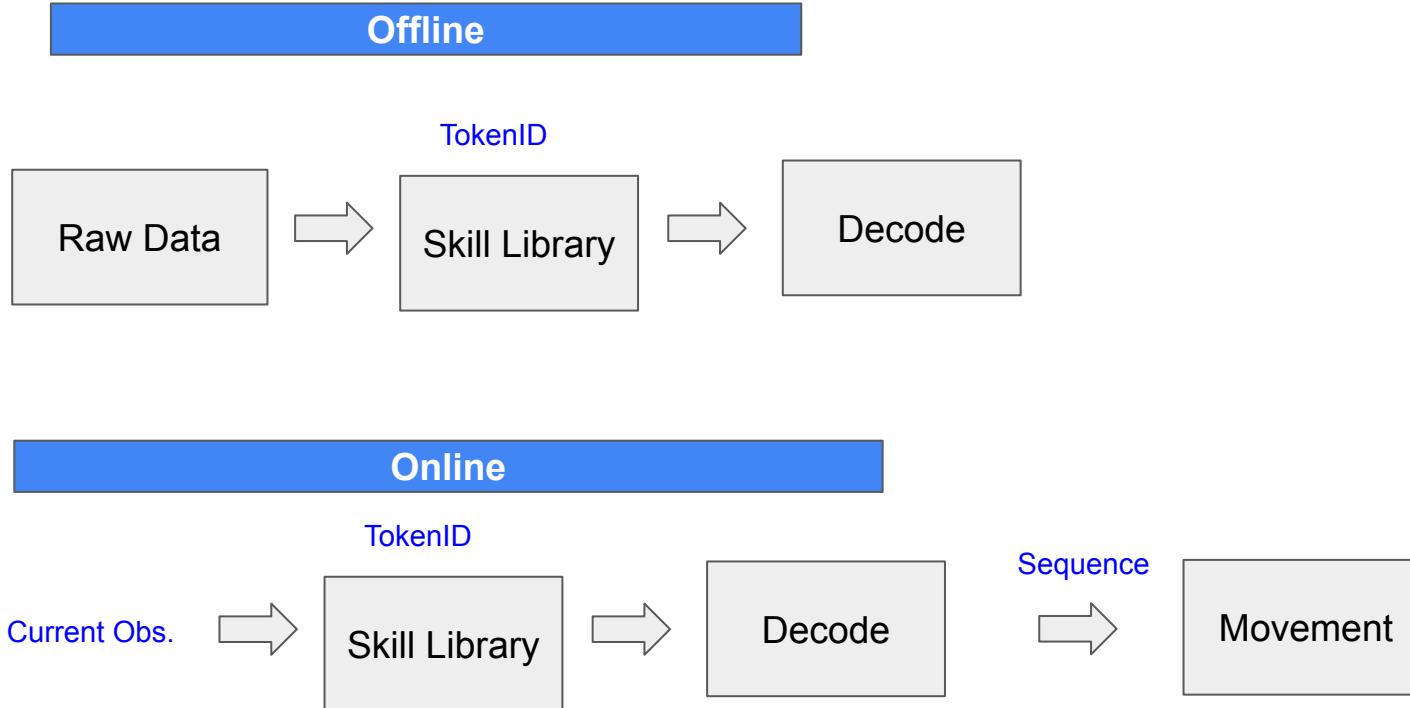
- If $\text{pooled_acc} \ll \text{oracle_acc}$ → each group has unique patterns (low composability)

Pooled: "I learned from all Romance languages, they share grammar!"
→ Translates Spanish well (90% accuracy)

Oracle: "I only learned Spanish"
→ Translates Spanish slightly better (92% accuracy)

Ratio: $90/92 = 0.98$ <- Good Composability

Overall Pipeline



Detail Example

[VAE Pretraining]

"Learn a good dictionary of action patterns"

Action [0.1, 0.2, 0.3, ...] → "code #15"

